

Jyotirmaya Sahu

AI Engineer | LLMs & Generative AI | RAG & Agentic Systems

+971 543334174 sahu.jyotirmaya26@gmail.com linkedin.com/in/jyotirmaya-sahu jyoti-dev-portfolio.vercel.app Dubai, UAE

EXPERIENCE

AI Engineer, Mindfire Solutions

Client: BrightEdge (enterprise SEO & content platform) · 5+ years

- Worked for 5+ years as BrightEdge's engineer on their AI product line, building the LLM features and AI analytics used across their enterprise customers.

AI PROJECTS

Content Advisor, AI Content Generation

LLM engine for long-form content

- Built BrightEdge's first AI product, an LLM engine that writes optimized long-form content using the customer's own data, competitor pages, and site content.
- Designed the prompt and context pipeline that feeds crawl data, keyword signals, and competitor benchmarks into the model so output stays factual and on topic.
- Streamed model tokens live over WebSockets, which cut the perceived wait time in half.
- Added a versioned editor so teams can edit and compare drafts, plus PDF and Excel export.

Content Recommendations, AI Copilot

Generative suggestions over crawl data

- Integrated BrightEdge's AI Copilot to read large site crawls and generate suggestions for meta descriptions, titles, and H1/H2 tags.
- Wrapped the model output in a review workflow so specialists approve or edit before anything ships, which raised trust and task completion.
- Rebuilt the legacy PHP module in React with caching and debounced calls, improving user satisfaction by 40%.

Daily Pulse, AI Insights & Alerts

Automated SEO health and competitor tracking

- Built a platform that generates daily SEO insights and competitor signals, with email and in-app alerts to keep users informed.
- Reworked the data and insight pipeline to cut bugs by 40% and speed up development by 30%.

Dashboards & Platform

Analytics surfaces and delivery framework

- Fixed a serious bottleneck where 100+ charts took 2 to 3 minutes to load, bringing it under 5 seconds with lazy rendering and caching.
- Co-built a React microfrontend framework (Vite, PNPM, React 18) that let teams ship AI modules about 3x faster.
- Added Playwright headless rendering for clean PDF and PNG exports of analytics and AI output.

SUMMARY

AI Engineer with 5+ years building production AI products, including LLM content engines, AI copilots, and data insight platforms for large enterprise accounts.

Comfortable across the full LLM stack: Python, OpenAI and Anthropic APIs, prompt engineering, RAG with vector search, and agentic workflows, plus the streaming and review UX that make AI reliable in production.

Based in Dubai on a self-sponsored UAE residence visa, so I can join on short notice. Looking for AI Engineer roles with teams in the UAE.

KEY ACHIEVEMENTS

Shipped a first AI product

Lead engineer on BrightEdge's first LLM content generation product.

Star Performer, 3 months running

Recognized three months in a row for AI and frontend work.

95% faster load times

Cut load times up to 95% on enterprise AI analytics screens.

SKILLS

Languages

Python, TypeScript, JavaScript, SQL

LLMs & GenAI

OpenAI (GPT), Anthropic (Claude), Gemini, Llama, Mistral, transformers (BERT, GPT, Llama), tool calling, structured output, NLP

Prompting & Agents

Prompt engineering, context management, LangChain, LangGraph, LlamaIndex, Semantic Kernel, agentic workflows, guardrails

RAG & Retrieval

RAG, embeddings, semantic search, Pinecone, FAISS, Chroma, Qdrant, pgvector, chunking, re-ranking

ML & Deep Learning

PyTorch, TensorFlow, Keras, Hugging Face, scikit-learn, Pandas, NumPy, neural networks, computer vision, fine-tuning

MLOps & Cloud

AWS (SageMaker), Azure ML, GCP (Vertex AI), Databricks, Docker, Kubernetes, MLflow, Weights & Biases, Kubeflow, Airflow, Spark, CI/CD

Data & Backend

FastAPI, Flask, Node.js, REST, WebSockets, MySQL, MongoDB

Frontend & Tools

React, Next.js, React Query, MUI, Git, Playwright, Sentry, JIRA